

# NNDL: ICP3

Jyothi Kiran Boddeda  
700769023

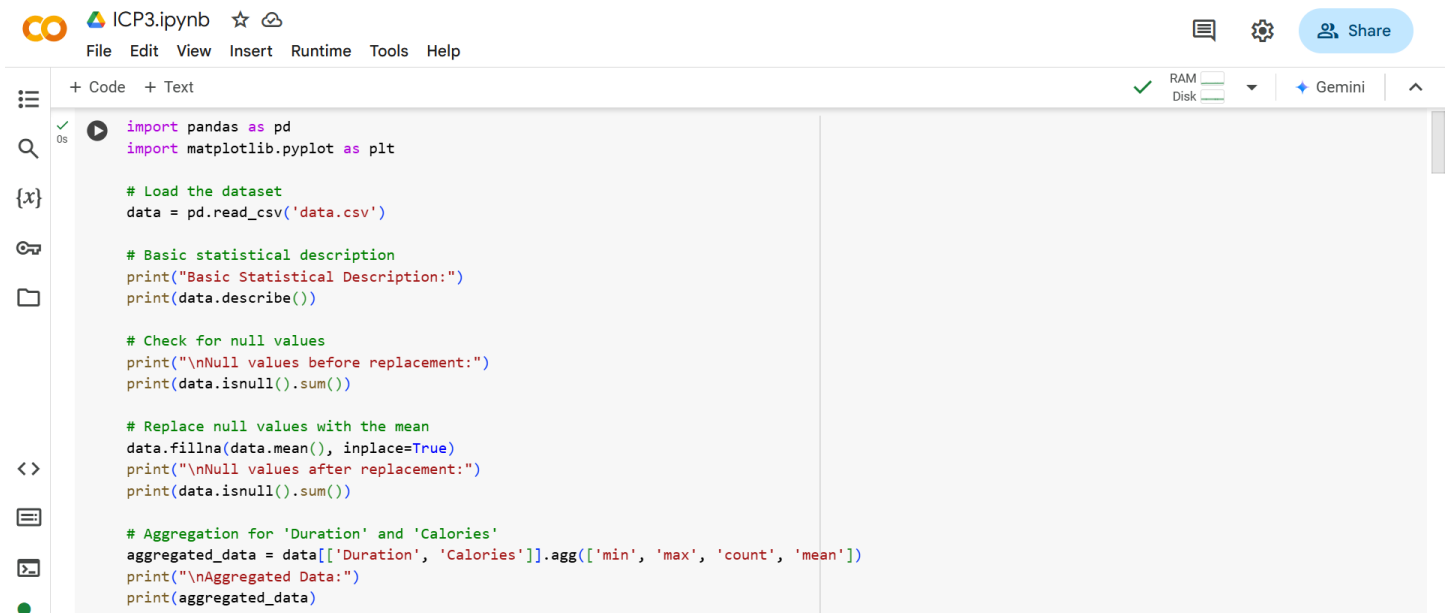
**GITHUB LINK :** <https://github.com/jyothikiranboddeda/Neural-Network-Deep-Learning.git>

**VIDEO LINK :**

[https://drive.google.com/file/d/1mxDSuX\\_u88JfdGzQjACyc6bl98Bi0I3B/view?usp=sharing](https://drive.google.com/file/d/1mxDSuX_u88JfdGzQjACyc6bl98Bi0I3B/view?usp=sharing)

## 1.Data Manipulation

- Read the provided CSV file 'data.csv'.
- <https://drive.google.com/drive/folders/1h8C3mLsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing>
- Show the basic statistical description about the data.
- Check if the data has null values.
- Replace the null values with the mean
- Select at least two columns and aggregate the data using: min, max, count, mean.
- Filter the dataframe to select the rows with calories values between 500 and 1000.
- Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
- Create a new "df\_modified" dataframe that contains all the columns from df except for "Maxpulse".
- Delete the "Maxpulse" column from the main df dataframe
- Convert the datatype of Calories column to int datatype.
- Using pandas create a scatter plot for the two columns (Duration and Calories).



```
ICP3.ipynb
File Edit View Insert Runtime Tools Help
+ Code + Text
import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
data = pd.read_csv('data.csv')

# Basic statistical description
print("Basic Statistical Description:")
print(data.describe())

# Check for null values
print("\nNull values before replacement:")
print(data.isnull().sum())

# Replace null values with the mean
data.fillna(data.mean(), inplace=True)
print("\nNull values after replacement:")
print(data.isnull().sum())

# Aggregation for 'Duration' and 'Calories'
aggregated_data = data[['Duration', 'Calories']].agg(['min', 'max', 'count', 'mean'])
print("\nAggregated Data:")
print(aggregated_data)
```



+ Code + Text

```
[10] # Filter rows with Calories between 500 and 1000
filtered_calories = data[(data['Calories'] >= 500) & (data['Calories'] <= 1000)]
print("\nRows with Calories between 500 and 1000:")
print(filtered_calories)

# Filter rows with Calories > 500 and Pulse < 100
filtered_pulse = data[(data['Calories'] > 500) & (data['Pulse'] < 100)]
print("\nRows with Calories > 500 and Pulse < 100:")
print(filtered_pulse)

# Create a new dataframe excluding the 'Maxpulse' column
df_modified = data.drop(columns=['Maxpulse'])
print("\nModified DataFrame without 'Maxpulse':")
print(df_modified.head())

# Delete 'Maxpulse' from the main dataframe
data.drop(columns=['Maxpulse'], inplace=True)

# Convert the Calories column to int
data['Calories'] = data['Calories'].astype(int)

# Scatter plot for 'Duration' and 'Calories'
plt.scatter(data['Duration'], data['Calories'], color='green')
plt.title('Scatter Plot of Duration vs Calories')
```



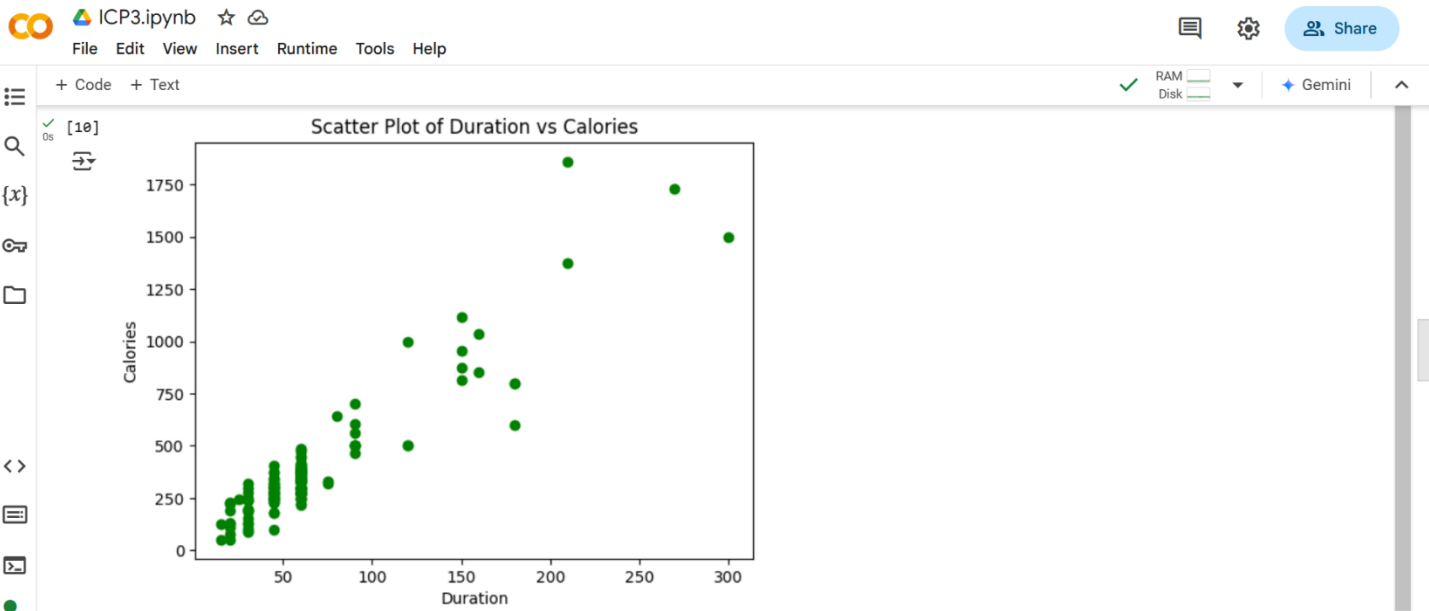
+ Code + Text

```
[10] # Scatter plot for 'Duration' and 'Calories'
plt.scatter(data['Duration'], data['Calories'], color='green')
plt.title('Scatter Plot of Duration vs Calories')
plt.xlabel('Duration')
plt.ylabel('Calories')
plt.show()
```

72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
83	120	100	130	500.0
90	180	101	127	600.1
99	90	93	124	604.1
101	90	90	110	500.0
102	90	90	100	500.0
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

Rows with Calories &gt; 500 and Pulse &lt; 100:

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2



## 2. Linear Regression

- Import the given “Salary\_Data.csv”
- Split the data in train\_test partitions, such that 1/3 of the data is reserved as test subset. c) Train and predict the model.
- Calculate the mean\_squared error
- Visualize both train and test data using scatter plot



The image shows a Jupyter Notebook interface with the following code:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# Import the dataset
salary_data = pd.read_csv('Salary_Data.csv')

# Split the data into train and test partitions
X = salary_data[['YearsExperience']]
y = salary_data['Salary']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1/3, random_state=42)

# Train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Calculate the Mean Squared Error
mse = mean_squared_error(y_test, y_pred)
```



The image shows a Jupyter Notebook interface with the following code:

```
# Calculate the Mean Squared Error
mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse}')

# Visualize the train and test data using scatter plot
plt.figure(figsize=(8, 4))

# Scatter plot for the training data
plt.scatter(X_train, y_train, color='blue', label='Train Data')

# Scatter plot for the test data
plt.scatter(X_test, y_test, color='red', label='Test Data')

# Plot the regression line
plt.plot(X_test, y_pred, color='black', linewidth=2, label='Regression Line')

plt.title('Salary vs Years of Experience')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend()
plt.show()
```

Mean Squared Error: 35301898.887134895

